Predict the Breach, Not the Drop: Control-Aligned Learning for Smoothing AI Training Power Swings

Sayed Raheel
sayedraheel1995@gmail.com

Abstract

Large AI training jobs move GPU power by tens of percent within a second, and synchronized jobs can violate the ramp-rate limits that grid operators place on large loads. Fill-type smoothing (filler GPU work or storage that adds power while the load falls) can remove these violations, but it has to be engaged before the fall arrives, and keeping it engaged costs energy. A forecaster can decide when to engage it. We show that what the forecaster is trained to predict matters more than which model is used. A gradient-boosted classifier trained on the usual target, "a large power drop is coming", ranks the events a ramp limiter must catch poorly (average precision, AP, 0.58 on held-out H100 traces). Retrained on the controller's own event, "a ramp-limit violation will occur while a limiter triggered now would be engaged", the same model class reaches AP 0.96, against 0.67 for a no-learning rhythm-replay baseline of the same target. In a limiter simulation with a threshold chosen on validation data for a 90% removal target, the learned trigger removes 86.8% of test violations with the actuator engaged 45.6% of the time, while rhythm replay meets the target only by staying engaged 99.2% of the time, i.e. effectively always on. Across three public datasets (H100, B200 and V100 nodes; benchmarks and real user jobs), the learned trigger ranks violations significantly better than rhythm replay on two datasets when the site is represented in training; the third is inconclusive. On a site never seen in training, the learned trigger is never better than rhythm replay in AP, and an hour of local data does not reliably recover the advantage. We report this as an open problem.

1 Introduction

Training a large neural network alternates compute-heavy and communication-heavy phases, so GPU power moves in fast, large steps. All nodes of a data-parallel job step together, so a synchronized job looks to the grid like one large, fast swing. Industry reports describe swings of tens to hundreds of megawatts and evaluate software power burning, GPU power floors and rack-level storage as mitigations [1]. Hardware filters and storage that react to the measured load are being developed [2, 3]; predictive storage control has been studied with perfect preview [4] and with short delay-compensating predictors [5]; and AI data center load has been forecast at minute scales [6].

This paper studies one specific decision: *when to engage* a fill-type actuator. The actuator (filler GPU kernels, a battery, a supercapacitor) adds power only while the load falls faster than the allowed ramp. It must be ready before the fall: kernels must be launched, an inverter must be hot. Readiness has a cost, so it should be engaged only when needed. A forecaster can make that call. The obvious competitor is free: training power is periodic, so "replay the last cycle" is a strong predictor.

Our contributions are:

- **A control-aligned prediction target.** Instead of a generic "large drop", the label is the event the controller acts on: a ramp-limit violation whose whole 1 s ramp window lies inside

Table 1: Datasets. Node-hours are those used for modeling. Test units are independent jobs or sessions.

Dataset	GPUs/node	Workloads	Sensor	Node-h	Units (test)
NLR [7]	4×H100	MLPerf Llama2-70B LoRA, SD; vLLM inference	NVML 5–10 Hz	248	98 (20)
Elsayed et al. [8]	8×H100 / B200	diffusion, Llama 1B–8B fine-tuning	NVML	8	32 (6)
MIT Supercloud [9]	1–2×V100	7 model classes, real user jobs	nvidia-smi 10 Hz	1,292	560 (112)

the window during which a limiter triggered now would be armed. Retargeting alone raises AP from 0.58 to 0.96 on held-out H100 traces (Section 4.1).

- **A readiness analysis** that puts learned triggers, rhythm replay, an always-on limiter and a perfect-foresight oracle on one curve of violations removed versus time engaged, with operating points chosen on validation data and scored on test data (Section 4.2).
- **A three-dataset study** (H100, B200, V100; scripted benchmarks and real user jobs) with cluster-bootstrap confidence intervals, a sensor-harmonization step, and a **negative result**: a learned trigger did not transfer to an unseen site (Sections 4.3–4.4).

We are not aware of prior work that studies how the forecasting target affects smoothing performance.

2 Data

Table 1 summarizes the three public datasets.

NLR [7] records 4×H100 nodes running MLPerf Llama2-70B LoRA and Stable Diffusion training on 1 to 16 nodes, plus three vLLM Llama3.1-70B inference benchmarks (offline, online with a finite request set, online at a fixed rate). We keep at most four nodes per training job; these run nearly in lockstep and all are weighted equally in the metrics. Training jobs are split whole into train, validation and test. Each inference benchmark is a single continuous session, cut into one-hour blocks and split in time; its test blocks therefore count as a single independent cluster, which gives 12 test clusters in total.

Elsayed et al. [8] record single 8×H100 and 8×B200 nodes running diffusion training and Llama 1B–8B fine-tuning sweeps in 15-minute sessions (32 unique sessions).

MIT Supercloud [9]: we use a stratified subset of 560 labelled jobs, 80 per model class (VGG, ResNet, Inception, U-Net, BERT, DistilBERT, GNN), on nodes with one or two V100 GPUs. These are real user jobs, not benchmarks.

Common preprocessing. All traces are averaged onto a 0.2 s grid (row t averages $[t, t+0.2\,\mathrm{s})$) and expressed as "NLR-equivalent watts": node GPU power times 2.8 kW divided by the node's rated GPU power. A ramp limit of 10% of rated power per second therefore means the same thing on every node. Jobs are assigned whole to train, validation and test (about 60/20/20, seed 42). For the control-aligned models the validation units are split again into two halves: one for early stopping and one for choosing operating points. For NLR inference these halves are alternate hour blocks of the same session and are not independent.

Sensor harmonization. NVML documents that its default power reading is averaged over a 1 s interval on Ampere (except GA100) and newer GPUs, which includes the H100 and B200, while on GA100 and older architectures, including the V100, it returns instantaneous power [10]. The NLR documentation does not state which reading was logged. Raw V100 traces flicker: without

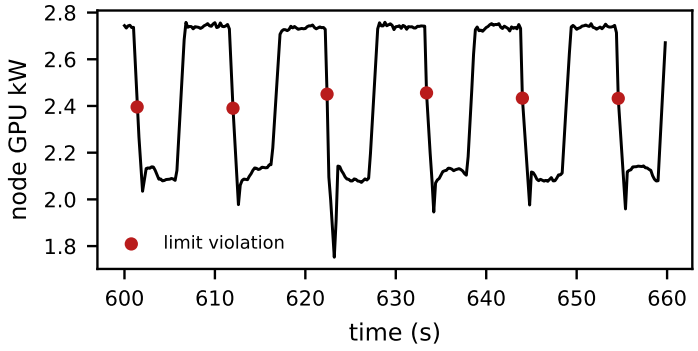


Figure 1: Thirty seconds of one NLR Llama2-70B LoRA test node. Red dots mark ramp-limit violations (1 s fall larger than 10% of rated node power). Many are steps of a regular square wave rather than "big drops".

harmonization, 89% of Supercloud seconds would be labelled "violation imminent". We apply a 1 s trailing mean to Supercloud so that all three datasets are compared at a similar effective bandwidth. This is an assumption, discussed in Section 5.

3 Problem and methods

3.1 Ramp limit, limiter and readiness

Let p_t be node power at row t (0.2 s steps) and $F = 5$ the number of rows per second. A down-ramp violation starts at row v when $p_{v-F} - p_v > 0.1\,P_{\text{rated}}$ and this did not hold at $v-1$; its ramp window is $[v - F, v]$ (Fig. 1).

The *limiter* adds power only. When armed, or while still filling, its output is

$$o_t = \max\big(p_t,\ o_{t-1} - 0.1\,P_{\text{rated}}/F\big), \tag{1}$$

so it tapers at the allowed rate until it meets the load; otherwise $o_t = p_t$.

A trigger issued at row t acts after a one-row decision delay (row t is complete only at $t + 1$) and an actuator latency L (1 s unless stated), so engagement starts at $e = t + 1 + LF$. Each trigger arms the limiter for the half-open window $[e, e + H)$ with hold $H = 6$ s $= 30$ rows. We count as *engaged* every row where the limiter is armed or still filling. Performance is the share of raw violations removed against the share of time engaged. Energy overhead is the fill energy plus a standby draw of 0 to 5% of rated power while engaged, as a share of job energy.

References: *always-on* (always armed), *always-burn* (an 80% power floor, which removes only 6.7% of violations for 6.8% extra energy), and an *oracle* that arms just in time for every violation, iterating until violations induced by the limiter itself are also covered.

3.2 Targets

The **big-drop** target labels row t positive if a fall of more than 30% of the trailing 10 s peak within 1 s starts in the next 5 s. The **control-aligned** target labels row t positive if a raw violation starts at some v with

$$v \in [e + F,\ e + H - 1], \qquad e = t + 1 + LF, \tag{2}$$

that is, the whole ramp window $[v-F, v]$ lies inside the armed window $[e, e+H)$. Violations induced by the limiter itself are not labelled. Many violations are not big drops: the Llama LoRA workload is a square wave of about 22% of rated power (Fig. 1), so the two targets differ substantially. We train one model per actuator latency $L \in \{0, 0.2, 0.5, 1, 2, 3\}$ s.

3.3 Predictors

Rhythm replay estimates the period P from a trailing 120 s autocorrelation (recomputed every second), predicts $\hat{p}_{t+k} = p_{t+k-P\lceil k/P \rceil}$, and scores row t by the largest predicted 1 s down-ramp whose window lies inside the same target window. It needs no training.

The **learned** predictor is LightGBM [11] on 56 trailing features: current power, lags and differences over 0.2 to 20 s, rolling mean, standard deviation, minimum and maximum over 1 to 60 s, time since the last up-step, down-step and drop, autocorrelation and FFT rhythm estimates (period, strength, phase) and rhythm-replay forecasts at 1, 5 and 10 s. It uses learning rate 0.05, 63 leaves, at least 200 rows per leaf, square-root class weighting and early stopping on the early-stopping half of validation. A **shape-only** variant, used for transfer, divides every level feature by the trailing 60 s maximum.

3.4 Evaluation

Ranking quality is AP on test rows. Uncertainty is a paired bootstrap over independent clusters (training jobs; each inference session is one cluster; 1000 resamples). The intervals resample test clusters for a single fitted model and split; they do not retrain, so they omit training and split variability. Operating points (thresholds) are chosen on the threshold half of validation and then scored on test. For the readiness curves we sweep about 200 thresholds taken from validation-score quantiles. Cross-dataset comparisons score one decision per second; the NLR readiness analysis scores every 0.2 s row.

4 Results

4.1 The training target

On NLR test traces rhythm replay ranks the control-aligned event with AP 0.67. LightGBM trained on the big-drop target reaches only 0.58 on the same event. The same model class trained on the control-aligned target reaches 0.96, and the gain holds for every workload (Table 2). The inference workloads are hardest; their falls are dominated by the gaps between benchmark runs. Results barely depend on actuator latency (Table 3): the regular structure that makes violations predictable 1 s ahead also makes them predictable 3 s ahead.

4.2 Readiness on NLR

Fig. 2 sweeps each trigger's threshold on NLR test traces. Every trigger can reach high removal by staying engaged long enough; what differs is how much engaged time it needs.

Validation-chosen operating points are the deployable result (Table 4). For a 90% removal target, the learned control-aligned trigger removes 86.8% of test violations at 45.6% engaged time, slightly short of target. Rhythm replay reaches 99.8% only by staying engaged 99.2% of the time, which is the always-on solution. With a 2% standby draw the energy overheads are 3.7% and 5.4% (always-on: 5.4%). At an 80% target both triggers undershoot on test (73.7% and 74.4%):

Table 2: NLR test AP on the control-aligned target, 1 s latency.

Workload	Learned	Rhythm replay
Llama2-70B LoRA training	1.00	0.81
Stable Diffusion training	0.96	0.73
Offline inference	0.55	0.29
Online inference, finite	0.76	0.09
Online inference, fixed rate	0.65	0.36
All	0.96	0.67

Table 3: NLR test AP by actuator latency.

L (s)	Learned	Rhythm
0	0.96	0.68
0.2	0.96	0.68
0.5	0.96	0.68
1	0.96	0.67
2	0.96	0.67
3	0.95	0.67

Table 4: Validation-chosen operating points scored on NLR test traces (1 s latency). Removal and engaged time in % ; overhead in % of job energy with a 2% standby draw.

Target	Trigger	Removed	Engaged	Overhead
80%	Learned, aligned	73.7	35.9	2.9
	Rhythm, aligned	74.4	43.3	3.2
	Learned, big-drop	88.0	61.4	4.3
90%	Learned, aligned	86.8	45.6	3.7
	Rhythm, aligned	99.8	99.2	5.4
	Learned, big-drop	92.7	69.1	4.6
95%	Learned, aligned	91.6	49.7	4.0
	Rhythm, aligned	99.8	99.2	5.4
	Learned, big-drop	95.4	75.8	4.8
	Always-on	100.0	100.0	5.4
	Oracle	100.0	29.8	3.6

thresholds chosen on a few validation jobs transfer imperfectly. The oracle removes all violations at 29.8% engaged time, which bounds what better prediction could still gain.

Attainable trade-offs (retrospective). Table 5 reads, from each test curve, the least engaged time at which a removal level is reached. Because the threshold is picked on test, this compares attainable trade-offs rather than deployable operating points. To remove 90% of violations the learned control-aligned trigger needs 48.4% engaged time and rhythm replay 99.2%: a paired cluster-bootstrap difference of −50.8 points (95% CI −57.4 to −8.8). The interval is asymmetric because the statistic is a minimum over thresholds: in most resamples rhythm replay needs more than 95% engagement, but some cluster mixtures cross 90% removal near 50–60%. At 95% removal the difference is −40.8 points (CI −49.1 to −31.4). Against the learned big-drop trigger the difference is −15.9 points at 90% (CI −25.7 to −3.9) but not significant at 95% (CI −26.3 to 0.5); with 12 test clusters the effect of the target is clearest in ranking quality.

4.3 Across sites and GPUs

Table 6 and Fig. 3 compare AP in three configurations: trained on the test site's own training jobs (in-domain), trained on all three datasets (pooled), and trained on the other two only (unseen site). In-domain or pooled, the learned trigger ranks violations significantly better than rhythm replay on NLR and on real Supercloud jobs. On Supercloud, exploratory per-class intervals (not corrected for multiplicity) are positive for 7 of 7 model classes. On Elsayed the point estimate favours learning

Table 5: Retrospective: least engaged time (%) on the NLR test curve that reaches each removal level, 1 s latency. Thresholds are read off the test curve.

Removal	Learned, aligned	Rhythm, aligned	Learned, big-drop
80%	43.1	51.5	53.0
90%	48.4	99.2	64.3
95%	58.4	99.2	74.8

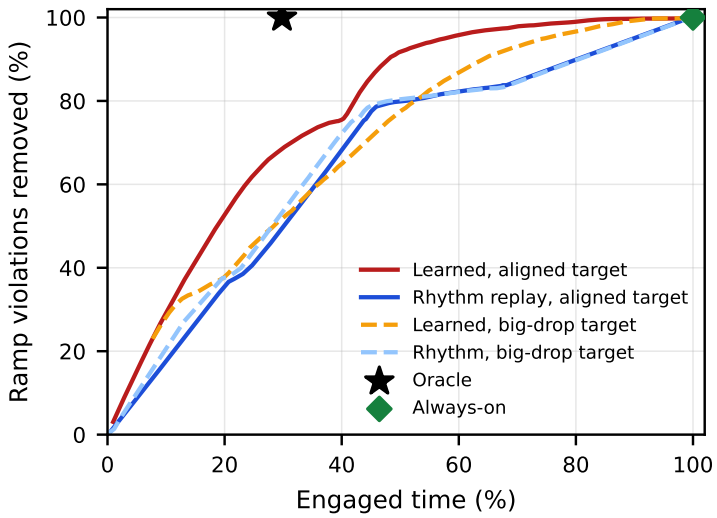


Figure 2: Violations removed versus engaged time on NLR test traces (1 s latency, about 200 thresholds from validation-score quantiles). Always-on removes everything but is always engaged; the oracle bounds what any trigger could achieve.

but the interval includes zero with only six test sessions, and rhythm replay is better on its two diffusion test sessions; their very regular cadence may suit replay, but two sessions do not allow a firm conclusion.

AP is a ranking measure, not a smoothing result. Table 7 reports validation-chosen operating points for a 90% removal target on each dataset (one decision per second, no confidence intervals). The learned trigger lands below target on Elsayed and slightly below on the others; rhythm replay meets the target on NLR and Supercloud only by staying engaged almost all the time, and misses it on Elsayed. Operationally, the main advantage of learning is engaged time, not removal.

4.4 Unseen sites

When the test dataset is excluded from training, the learned trigger is not better than rhythm replay in AP on any dataset, and with the ordinary features it is significantly worse on NLR and Supercloud (Table 6). Absolute power levels differ across GPU generations even after normalization, so we tried two remedies (Table 8). Shape-only features close the NLR gap to an inconclusive difference but remain significantly worse on Supercloud. Continuing to boost the unseen-site model on 5, 15 or 60 minutes of the new site's training data (mean and range over five independent draws) does not reliably recover the in-domain advantage. Rhythm replay needs no site data at all; in our experiments a learned trigger needed the site in its training set.

Validation-chosen operating points tell a slightly different story: with ordinary features, the

Table 6: Test AP, learned versus rhythm replay, with the 95% cluster-bootstrap interval of the difference. Unseen site uses the ordinary features (shape-only results in Table 8).

Test dataset	Training data	Learned	Rhythm	Difference CI
NLR	in-domain	0.96	0.67	[0.22, 0.36]
	pooled	0.96	0.67	[0.22, 0.37]
	unseen site	0.53	0.67	[−0.23, −0.02]
Supercloud	in-domain	0.88	0.74	[0.11, 0.19]
	pooled	0.88	0.74	[0.11, 0.18]
	unseen site	0.68	0.74	[−0.10, −0.02]
Elsayed	in-domain	0.93	0.88	[−0.11, 0.09]
	pooled	0.92	0.88	[−0.13, 0.08]
	unseen site	0.84	0.88	[−0.43, 0.05]

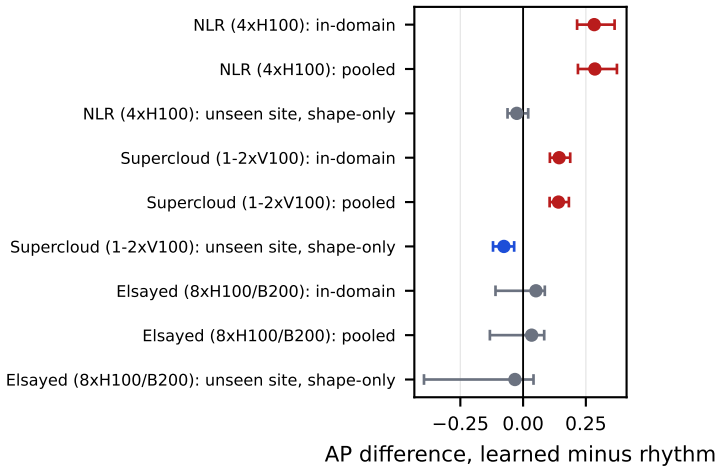


Figure 3: AP difference, learned control-aligned trigger minus rhythm replay, with 95% cluster-bootstrap intervals. Red: learned significantly better; blue: rhythm significantly better; grey: inconclusive. The unseen-site rows use shape-only features.

unseen-site learned trigger on NLR removes 90.3% of violations at 58.0% engaged time, against 99.8% at 99.1% for rhythm replay. Ranking quality and engaged time can therefore disagree; we have no interval for this comparison and do not claim transfer from it.

5 Discussion and limitations

Why the target matters. The big-drop label teaches the model about one kind of event (large falls) at one horizon (the next 5 s). The limiter needs something else: a fall steep enough to break the ramp limit, inside the window its own latency and hold time define. Many violations, such as the steps of the LoRA square wave, never qualify as big drops, so a big-drop model never learns to rank them. The model class, features and data were the same in both cases; only the label changed.

Why rhythm replay is always on at high removal. Replay predicts each violation from the previous cycle. Violations that do not repeat (phase changes, gaps between inference runs) score near zero, so the only way for replay to catch them is to lower the threshold until the limiter is armed almost all the time.

Table 7: Validation-chosen operating points for a 90% removal target, scored on test (pooled training; one decision per second). Removed / engaged, in %.

Test dataset	Learned	Rhythm replay
NLR	88.2 / 44.0	99.8 / 99.1
Supercloud	87.7 / 53.6	99.8 / 99.3
Elsayed	83.4 / 55.4	67.3 / 51.8

Table 8: Transfer to an unseen site with shape-only features. Adapted rows: mean (range) AP over five draws of local data. Rhythm replay AP for reference.

	NLR	Supercloud	Elsayed
Rhythm replay	0.67	0.74	0.88
In-domain, shape-only	0.96	0.88	0.88
Unseen site, shape-only	0.65	0.66	0.85
difference CI	[−0.06, 0.02]	[−0.12, −0.04]	[−0.39, 0.04]
+ 5 min of site data	0.63 (0.56–0.65)	0.62 (0.59–0.66)	0.82 (0.76–0.85)
+ 15 min	0.59 (0.55–0.65)	0.64 (0.57–0.67)	0.78 (0.66–0.88)
+ 60 min	0.71 (0.62–0.77)	0.64 (0.58–0.69)	0.84 (0.80–0.91)

Limitations. *Sensors.* If NLR and Elsayed power is NVML's 1 s average, true ramps are faster than measured and measured lead times are optimistic; harmonizing Supercloud with a 1 s mean is an assumption. *Uncertainty.* Intervals are conditional on one fitted model and one split and do not include training or split variability; NLR has 12 test clusters and Elsayed six. Up to four near-synchronous nodes per NLR job are weighted equally, so larger jobs weigh more in point estimates. For NLR inference, the two validation halves are not independent. *Simulation.* The actuator is an idealized limiter with fixed latency, hold time and standby draw. Storage capacity, state of charge, recharge, efficiency, wear and measured actuator dynamics are not modelled, and the standby cost is an assumed fraction of rated power. *Workloads.* NLR and Elsayed are scripted benchmarks; NLR inference falls are dominated by gaps between benchmark runs, and each NLR inference workload is a single session. Supercloud adds real jobs, but on 2018-era V100 nodes. *Scale.* All results are per node. Synchronized multi-node jobs scale the same swing, while desynchronized jobs partly cancel; we do not model facility-level aggregation.

Future work. Repeated whole-job splits or a nested bootstrap that retrains; one-node-per-job sensitivity; a storage model with state of charge; and methods that let a learned trigger transfer to a new site, or at least decide quickly whether to fall back to rhythm replay.

6 Conclusion

For fill-type smoothing of AI power swings, a forecaster should be trained on the controller's event, a ramp-limit violation inside the actuator's engagement window, rather than on a generic large drop. With that target, a small gradient-boosted model needs substantially less engaged time than rhythm replay for the same violation removal on H100 benchmark traces, and ranks violations better (higher AP) on real V100 jobs, provided the site is represented in training. Closing the gap to the oracle and making learned triggers transfer to new sites are the natural next steps.

Data, code and licences

The NLR data are released under the National Laboratory of the Rockies dataset licence, which permits use with attribution. The Elsayed et al. dataset is listed as CC BY 4.0 on figshare, while the README inside its archive states CC BY-NC-ND; we treat it conservatively as non-commercial. The MIT Supercloud data are CC BY-NC-ND 4.0. All three are used for research only and none is redistributed; readers can obtain them from the original sources. Every number in this paper is generated from result tables written by the analysis pipeline. The code is available from the author on request.

Use of AI tools

Analysis code and text were drafted with the help of AI coding and writing assistants and independently reviewed by a separate AI model. The author checked the results and is responsible for the content.

References

[1] E. Choukse, B. Warrier, S. Heath, et al., "Power stabilization for AI training datacenters," arXiv:2508.14318, 2025.

[2] D. Jensen, G. Wilkins, O. Nnorom Jr., H. Budd, R. Rajagopal, J. Rivas-Davila, and P. Levis, "EasyRider: Mitigating power transients in datacenter-scale training workloads," arXiv:2604.15522, 2026.

[3] D. Jensen, G. Wilkins, O. Nnorom Jr., H. Budd, P. Levis, and J. Rivas, "Active filter design for buffering datacenter-scale power fluctuations from training AI," arXiv:2609.13490, 2026.

[4] H. You, C. Lou, and J. Yang, "Source side mitigation of AI datacenter power fluctuations with a hybrid energy storage system and residual differentiable predictive control," arXiv:2606.04869, 2026.

[5] X. Lyu, W. Du, S. M. Mohiuddin, and B. A. Ross, "Communication delay robust control of BESS for AI training load smoothing," arXiv:2609.22542, 2026.

[6] M. Mughees, Y. Li, Y. Chen, and Y. R. Li, "Short-term load forecasting for AI-data center," arXiv:2503.07756, 2025 (accepted to IEEE PES General Meeting 2025).

[7] R. Vercellino, J. Willard, G. Campos, W. da Silva Pereira, O. Hull, M. Selensky, and J. Mueller, "Measurement of generative AI workload power profiles for whole-facility data center infrastructure planning," arXiv:2604.07345, 2026. Data: "Dataset of Generative AI Workload Power Profiles," National Laboratory of the Rockies, doi:10.7799/3025227.

[8] A. A. E. Elsayed, A. A. Al-Obaidi, and H. E. Z. Farag, "Characterization of high-resolution AI data center training workloads on single and multiple GPU nodes," *Scientific Data*, vol. 13, art. 1268, 2026, doi:10.1038/s41597-026-07496-6. Data: figshare, doi:10.6084/m9.figshare.31654879.

[9] S. Samsi, M. L. Weiss, D. Bestor, et al., "The MIT Supercloud dataset," in *Proc. IEEE High Performance Extreme Computing Conf. (HPEC)*, 2021, pp. 1–8, doi:10.1109/HPEC49654.2021.9622850.

[10] NVIDIA, "NVML API reference: nvmlDeviceGetPowerUsage," https://docs.nvidia.com/deploy/nvml-api/api/group__nvmlDeviceQueries.html, accessed 2026-10-15.

[11] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.